

```

import torch
torch.manual_seed(17)

import numpy as np
from torchsummary import summary
from tqdm import tqdm
import matplotlib.pyplot as plt

from DatasetLoader import DatasetFetcher
from project_model import *

# if torch.backends.mps.is_available():
#     mps_device = torch.device("mps")
#     x = torch.ones(1, device=mps_device)
#     print (x)
# else:
#     print ("MPS device not found.")

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

cuda

# Fetching Dataset
df = DatasetFetcher(dataset = "CIFAR10", batch_size = 128)
df.addHorizontalFlipping()
#df.addVerticalFlipping()
df.addRandomCrop(size = 32, padding = 4)
#df.addAutoAugmentation()
#df.addHistogramEqualization()
df.addNormalizer()
#df.addGaussianNoise()
trainLoader, testLoader = df.getLoaders()

Initializing fetching CIFAR10 dataset using torchvision
Files already downloaded and verified
Files already downloaded and verified
Files already downloaded and verified

# Get Model
#model = ResNet(BasicBlock, 32, 4, [4, 4, 4, 2], 10, bias=True)
model = project1_model()
model = model.to(device)
print(summary(model, input_size = (3, 32, 32)))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	896
BatchNorm2d-2	[-1, 32, 32, 32]	64
Conv2d-3	[-1, 32, 32, 32]	9,248

BatchNorm2d-4	[-1, 32, 32, 32]	64
Conv2d-5	[-1, 32, 32, 32]	9,248
BatchNorm2d-6	[-1, 32, 32, 32]	64
BasicBlock-7	[-1, 32, 32, 32]	0
Conv2d-8	[-1, 32, 32, 32]	9,248
BatchNorm2d-9	[-1, 32, 32, 32]	64
Conv2d-10	[-1, 32, 32, 32]	9,248
BatchNorm2d-11	[-1, 32, 32, 32]	64
BasicBlock-12	[-1, 32, 32, 32]	0
Conv2d-13	[-1, 32, 32, 32]	9,248
BatchNorm2d-14	[-1, 32, 32, 32]	64
Conv2d-15	[-1, 32, 32, 32]	9,248
BatchNorm2d-16	[-1, 32, 32, 32]	64
BasicBlock-17	[-1, 32, 32, 32]	0
Conv2d-18	[-1, 32, 32, 32]	9,248
BatchNorm2d-19	[-1, 32, 32, 32]	64
Conv2d-20	[-1, 32, 32, 32]	9,248
BatchNorm2d-21	[-1, 32, 32, 32]	64
BasicBlock-22	[-1, 32, 32, 32]	0
Conv2d-23	[-1, 64, 16, 16]	18,496
BatchNorm2d-24	[-1, 64, 16, 16]	128
Conv2d-25	[-1, 64, 16, 16]	36,928
BatchNorm2d-26	[-1, 64, 16, 16]	128
Conv2d-27	[-1, 64, 16, 16]	2,112
BatchNorm2d-28	[-1, 64, 16, 16]	128
BasicBlock-29	[-1, 64, 16, 16]	0
Conv2d-30	[-1, 64, 16, 16]	36,928
BatchNorm2d-31	[-1, 64, 16, 16]	128
Conv2d-32	[-1, 64, 16, 16]	36,928
BatchNorm2d-33	[-1, 64, 16, 16]	128
BasicBlock-34	[-1, 64, 16, 16]	0
Conv2d-35	[-1, 64, 16, 16]	36,928
BatchNorm2d-36	[-1, 64, 16, 16]	128
Conv2d-37	[-1, 64, 16, 16]	36,928
BatchNorm2d-38	[-1, 64, 16, 16]	128
BasicBlock-39	[-1, 64, 16, 16]	0
Conv2d-40	[-1, 64, 16, 16]	36,928
BatchNorm2d-41	[-1, 64, 16, 16]	128
Conv2d-42	[-1, 64, 16, 16]	36,928
BatchNorm2d-43	[-1, 64, 16, 16]	128
BasicBlock-44	[-1, 64, 16, 16]	0
Conv2d-45	[-1, 128, 8, 8]	73,856
BatchNorm2d-46	[-1, 128, 8, 8]	256
Conv2d-47	[-1, 128, 8, 8]	147,584
BatchNorm2d-48	[-1, 128, 8, 8]	256
Conv2d-49	[-1, 128, 8, 8]	8,320
BatchNorm2d-50	[-1, 128, 8, 8]	256
BasicBlock-51	[-1, 128, 8, 8]	0
Conv2d-52	[-1, 128, 8, 8]	147,584
BatchNorm2d-53	[-1, 128, 8, 8]	256

Conv2d-54	[-1, 128, 8, 8]	147,584
BatchNorm2d-55	[-1, 128, 8, 8]	256
BasicBlock-56	[-1, 128, 8, 8]	0
Conv2d-57	[-1, 128, 8, 8]	147,584
BatchNorm2d-58	[-1, 128, 8, 8]	256
Conv2d-59	[-1, 128, 8, 8]	147,584
BatchNorm2d-60	[-1, 128, 8, 8]	256
BasicBlock-61	[-1, 128, 8, 8]	0
Conv2d-62	[-1, 128, 8, 8]	147,584
BatchNorm2d-63	[-1, 128, 8, 8]	256
Conv2d-64	[-1, 128, 8, 8]	147,584
BatchNorm2d-65	[-1, 128, 8, 8]	256
BasicBlock-66	[-1, 128, 8, 8]	0
Conv2d-67	[-1, 256, 4, 4]	295,168
BatchNorm2d-68	[-1, 256, 4, 4]	512
Conv2d-69	[-1, 256, 4, 4]	590,080
BatchNorm2d-70	[-1, 256, 4, 4]	512
Conv2d-71	[-1, 256, 4, 4]	33,024
BatchNorm2d-72	[-1, 256, 4, 4]	512
BasicBlock-73	[-1, 256, 4, 4]	0
Conv2d-74	[-1, 256, 4, 4]	590,080
BatchNorm2d-75	[-1, 256, 4, 4]	512
Conv2d-76	[-1, 256, 4, 4]	590,080
BatchNorm2d-77	[-1, 256, 4, 4]	512
BasicBlock-78	[-1, 256, 4, 4]	0
Linear-79	[-1, 10]	2,570

```
=====
Total params: 3,576,842
Trainable params: 3,576,842
Non-trainable params: 0
```

```
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 10.00
Params size (MB): 13.64
Estimated Total Size (MB): 23.66
-----
```

None

```
EPOCHS = 100
globalBestAccuracy = 0.0
trainingLoss = []
testingLoss = []
trainingAccuracy = []
testingAccuracy = []
```

```
# Defining Loss Function, Learning Rate, Weight Decay, Optimizer)
lossFunction = torch.nn.CrossEntropyLoss(reduction = 'sum')
learningRate = 0.1
weightDecay = 0.0001
# optimizer = torch.optim.Adam(model.parameters(), lr=learningRate,
weight_decay=weightDecay)
```

```

# optimizer = torch.optim.Adagrad(model.parameters(), lr=learningRate,
weight_decay=weightDecay)
optimizer = torch.optim.Adadelata(model.parameters(), lr =
learningRate, weight_decay = weightDecay)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer,
EPOCHS, eta_min = learningRate/10.0)
print(model.eval())
trainable_parameters = sum(p.numel() for p in model.parameters() if
p.requires_grad)
print("Total Trainable Parameters : %s"%(trainable_parameters))
if trainable_parameters > 5 * (10 ** 6):
    raise Exception("Model not under budget!")

ResNet(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
    (1): BasicBlock(
      (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
    (2): BasicBlock(
      (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (shortcut): Sequential()
    )
    (3): BasicBlock(
      (conv1): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential(
        (0): Conv2d(32, 64, kernel_size=(1, 1), stride=(2, 2))
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
    (2): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (shortcut): Sequential()
    )
    (3): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
  )
  (layer3): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2))
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
    (2): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (shortcut): Sequential()
    )
    (3): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
  )
  (layer4): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential(
        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2))
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (shortcut): Sequential()
    )
  )
  (linear): Linear(in_features=256, out_features=10, bias=True)
)
Total Trainable Parameters : 3576842

```

Training

```

for i in tqdm(range(EPOCHS)):
    for phase in ['train', 'test']:
        if phase == "train":

```

```

        loader = trainLoader
        model.train()
        optimizer.zero_grad()
    else:
        loader = testLoader
        model.eval()
    runningLoss = 0.0
    runningCorrects = 0
    for images, labels in loader:
        images = images.to(device)
        labels = labels.to(device)
        output = model(images)
        loss = lossFunction(output, labels)
        predicted_labels = torch.argmax(output, dim=1)
        #runningLoss += loss.item()*images.size(0)
        runningLoss += loss.item()
        runningCorrects += torch.sum(predicted_labels ==
labels).float().item()
        if phase == "train":
            loss.backward()
            optimizer.step()
    epochLoss = runningLoss/len(loader.dataset)
    epochAccuracy = runningCorrects/len(loader.dataset)
    if phase == "train":
        scheduler.step()
        trainingLoss.append(epochLoss)
        trainingAccuracy.append(epochAccuracy)
    else:
        testingLoss.append(epochLoss)
        testingAccuracy.append(epochAccuracy)
        if epochAccuracy > globalBestAccuracy:
            globalBestAccuracy = epochAccuracy
            model.saveToDisk()
    print("Training Loss : %s, Testing Loss : %s, Training Accuracy :
%s, Testing Accuracy : %s"\
        %(trainingLoss[-1], testingLoss[-1], trainingAccuracy[-1],
testingAccuracy[-1]))

1%|          | 1/100 [00:10<18:08, 10.99s/it]

Training Loss : 1.750534892578125, Testing Loss : 1.7342327341079713,
Training Accuracy : 0.34224, Testing Accuracy : 0.3896

2%||         | 2/100 [00:21<17:32, 10.74s/it]

Training Loss : 1.5082956085205077, Testing Loss : 1.4530678709030151,
Training Accuracy : 0.44076, Testing Accuracy : 0.4681

3%||         | 3/100 [00:32<17:10, 10.63s/it]

```


Training Loss : 1.306699574584961, Testing Loss : 1.224788269996643,
Training Accuracy : 0.52454, Testing Accuracy : 0.5583

4%|█ | 4/100 [00:42<17:04, 10.67s/it]

Training Loss : 1.1761730513000488, Testing Loss : 1.131978190612793,
Training Accuracy : 0.57418, Testing Accuracy : 0.5967

5%|█ | 5/100 [00:53<16:43, 10.56s/it]

Training Loss : 1.045868921508789, Testing Loss : 0.9739497625350952,
Training Accuracy : 0.62728, Testing Accuracy : 0.654

6%|█ | 6/100 [01:03<16:23, 10.47s/it]

Training Loss : 0.9365599842834472, Testing Loss : 0.9753861132621765,
Training Accuracy : 0.66576, Testing Accuracy : 0.6545

7%|█ | 7/100 [01:14<16:19, 10.53s/it]

Training Loss : 0.8710366773986816, Testing Loss : 0.9070924603462219,
Training Accuracy : 0.69146, Testing Accuracy : 0.6908

8%|█ | 8/100 [01:24<16:03, 10.47s/it]

Training Loss : 0.7868001916503906, Testing Loss : 0.8118785608291625,
Training Accuracy : 0.72428, Testing Accuracy : 0.72

9%|█ | 9/100 [01:34<15:53, 10.48s/it]

Training Loss : 0.7216525913238525, Testing Loss : 0.838003656578064,
Training Accuracy : 0.74744, Testing Accuracy : 0.7098

10%|█ | 10/100 [01:45<15:47, 10.53s/it]

Training Loss : 0.6740130103302002, Testing Loss : 0.7190186059951782,
Training Accuracy : 0.76368, Testing Accuracy : 0.7558

11%|█ | 11/100 [01:56<15:37, 10.53s/it]

Training Loss : 0.616966068649292, Testing Loss : 0.6758499299526215,
Training Accuracy : 0.78718, Testing Accuracy : 0.7718

12%|█ | 12/100 [02:06<15:27, 10.54s/it]

Training Loss : 0.5796961762237549, Testing Loss : 0.6749031675338745,
Training Accuracy : 0.79728, Testing Accuracy : 0.7725

13%|█ | 13/100 [02:17<15:16, 10.54s/it]

Training Loss : 0.5621770449066162, Testing Loss : 0.677848752117157,
Training Accuracy : 0.80552, Testing Accuracy : 0.7785

14%|█ | 14/100 [02:27<14:59, 10.46s/it]

Training Loss : 0.521658950881958, Testing Loss : 0.5816180925369263,
Training Accuracy : 0.82026, Testing Accuracy : 0.8048

15%|■ | 15/100 [02:37<14:45, 10.41s/it]

Training Loss : 0.5030528082275391, Testing Loss : 0.5683472782611847,
Training Accuracy : 0.8271, Testing Accuracy : 0.8063

16%|■ | 16/100 [02:48<14:36, 10.44s/it]

Training Loss : 0.4880191324996948, Testing Loss : 0.5641240973472595,
Training Accuracy : 0.83208, Testing Accuracy : 0.8082

17%|■ | 17/100 [02:58<14:26, 10.45s/it]

Training Loss : 0.46874917739868166, Testing Loss :
0.5573918341636658, Training Accuracy : 0.83866, Testing Accuracy :
0.8151

18%|■ | 18/100 [03:09<14:12, 10.40s/it]

Training Loss : 0.43556508716583253, Testing Loss :
0.5247749044418335, Training Accuracy : 0.84942, Testing Accuracy :
0.8232

19%|■ | 19/100 [03:19<14:03, 10.41s/it]

Training Loss : 0.4199140442276001, Testing Loss : 0.5646599299430847,
Training Accuracy : 0.85444, Testing Accuracy : 0.8075

20%|■ | 20/100 [03:29<13:46, 10.33s/it]

Training Loss : 0.41516722038269044, Testing Loss :
0.5037532205581665, Training Accuracy : 0.8576, Testing Accuracy :
0.8308

21%|■ | 21/100 [03:40<13:40, 10.39s/it]

Training Loss : 0.3896746640777588, Testing Loss : 0.5259988674163818,
Training Accuracy : 0.86682, Testing Accuracy : 0.8253

22%|■ | 22/100 [03:50<13:33, 10.43s/it]

Training Loss : 0.3784850244140625, Testing Loss : 0.4869466375350952,
Training Accuracy : 0.86954, Testing Accuracy : 0.8357

23%|■ | 23/100 [04:01<13:25, 10.47s/it]

Training Loss : 0.3628396271133423, Testing Loss : 0.4769480683326721,
Training Accuracy : 0.87546, Testing Accuracy : 0.8399

24%|■ | 24/100 [04:11<13:08, 10.38s/it]

Training Loss : 0.34634707462310793, Testing Loss :
0.49430470991134645, Training Accuracy : 0.88002, Testing Accuracy :
0.8376

25%|██████████ | 25/100 [04:21<13:02, 10.44s/it]

Training Loss : 0.3391695451354981, Testing Loss : 0.493526510810852,
Training Accuracy : 0.88278, Testing Accuracy : 0.8397

26%|██████████ | 26/100 [04:32<12:52, 10.44s/it]

Training Loss : 0.32523911186218263, Testing Loss :
0.43323381824493407, Training Accuracy : 0.88692, Testing Accuracy :
0.8571

27%|██████████ | 27/100 [04:42<12:41, 10.44s/it]

Training Loss : 0.3070455528640747, Testing Loss : 0.4389434522628784,
Training Accuracy : 0.89318, Testing Accuracy : 0.8532

28%|██████████ | 28/100 [04:53<12:34, 10.48s/it]

Training Loss : 0.296951802444458, Testing Loss : 0.48018632173538206,
Training Accuracy : 0.89616, Testing Accuracy : 0.8493

29%|██████████ | 29/100 [05:03<12:22, 10.46s/it]

Training Loss : 0.2894485390853882, Testing Loss :
0.45177939550876617, Training Accuracy : 0.89862, Testing Accuracy :
0.8567

30%|██████████ | 30/100 [05:14<12:14, 10.49s/it]

Training Loss : 0.2819375555610657, Testing Loss : 0.4342577423095703,
Training Accuracy : 0.90124, Testing Accuracy : 0.8612

31%|██████████ | 31/100 [05:24<11:57, 10.40s/it]

Training Loss : 0.2773608359146118, Testing Loss :
0.45364204959869386, Training Accuracy : 0.9037, Testing Accuracy :
0.8555

32%|██████████ | 32/100 [05:35<11:54, 10.51s/it]

Training Loss : 0.2649358284568787, Testing Loss :
0.41601660060882567, Training Accuracy : 0.908, Testing Accuracy :
0.866

33%|██████████ | 33/100 [05:45<11:32, 10.34s/it]

Training Loss : 0.25196005603790284, Testing Loss :
0.43029733076095583, Training Accuracy : 0.91294, Testing Accuracy :
0.8628

34%|██████ | 34/100 [05:55<11:21, 10.32s/it]

Training Loss : 0.24615072959899903, Testing Loss :
0.4192706132411957, Training Accuracy : 0.91334, Testing Accuracy :
0.8685

35%|██████ | 35/100 [06:06<11:16, 10.41s/it]

Training Loss : 0.22745964323043824, Testing Loss :
0.41750600509643554, Training Accuracy : 0.91998, Testing Accuracy :
0.8717

36%|██████ | 36/100 [06:17<11:15, 10.56s/it]

Training Loss : 0.23327418605804442, Testing Loss :
0.43516239695549014, Training Accuracy : 0.91852, Testing Accuracy :
0.864

37%|██████ | 37/100 [06:28<11:15, 10.72s/it]

Training Loss : 0.22252747444152832, Testing Loss :
0.40060316247940064, Training Accuracy : 0.92252, Testing Accuracy :
0.8794

38%|██████ | 38/100 [06:38<10:58, 10.63s/it]

Training Loss : 0.19907417359352111, Testing Loss :
0.41780849165916445, Training Accuracy : 0.93024, Testing Accuracy :
0.8778

39%|██████ | 39/100 [06:49<10:46, 10.59s/it]

Training Loss : 0.20170173211097717, Testing Loss :
0.43313117980957033, Training Accuracy : 0.9291, Testing Accuracy :
0.8717

40%|██████ | 40/100 [06:59<10:34, 10.57s/it]

Training Loss : 0.19560000583648682, Testing Loss :
0.4325757998466492, Training Accuracy : 0.93202, Testing Accuracy :
0.8761

41%|██████ | 41/100 [07:09<10:15, 10.43s/it]

Training Loss : 0.19260304227828978, Testing Loss :
0.4230586800336838, Training Accuracy : 0.93324, Testing Accuracy :
0.8795

42%|██████ | 42/100 [07:20<10:04, 10.43s/it]

Training Loss : 0.1884927225112915, Testing Loss :
0.44363087105751037, Training Accuracy : 0.934, Testing Accuracy :
0.8737

43%|██████ | 43/100 [07:30<09:56, 10.47s/it]

Training Loss : 0.18156796796798705, Testing Loss :
0.42046577625274656, Training Accuracy : 0.93464, Testing Accuracy :
0.8784

44%|██████ | 44/100 [07:41<09:47, 10.50s/it]

Training Loss : 0.1728495922756195, Testing Loss : 0.4485450621128082,
Training Accuracy : 0.9383, Testing Accuracy : 0.877

45%|██████ | 45/100 [07:51<09:39, 10.53s/it]

Training Loss : 0.16596370498657226, Testing Loss :
0.43388833293914797, Training Accuracy : 0.94092, Testing Accuracy :
0.8772

46%|██████ | 46/100 [08:02<09:31, 10.58s/it]

Training Loss : 0.16174998712539673, Testing Loss :
0.39738049416542054, Training Accuracy : 0.94316, Testing Accuracy :
0.887

47%|██████ | 47/100 [08:13<09:19, 10.56s/it]

Training Loss : 0.14697849188804626, Testing Loss :
0.41343463668823244, Training Accuracy : 0.9477, Testing Accuracy :
0.8905

48%|██████ | 48/100 [08:23<09:06, 10.50s/it]

Training Loss : 0.13773472970962525, Testing Loss :
0.4124880380630493, Training Accuracy : 0.95168, Testing Accuracy :
0.8842

49%|██████ | 49/100 [08:33<08:49, 10.38s/it]

Training Loss : 0.1393139978981018, Testing Loss :
0.45448302216529846, Training Accuracy : 0.95046, Testing Accuracy :
0.884

50%|██████ | 50/100 [08:44<08:41, 10.43s/it]

Training Loss : 0.13646688494682313, Testing Loss :
0.4219742176771164, Training Accuracy : 0.95198, Testing Accuracy :
0.8878

51%|██████ | 51/100 [08:54<08:31, 10.44s/it]

Training Loss : 0.1224304889678955, Testing Loss : 0.4543603331565857,
Training Accuracy : 0.9567, Testing Accuracy : 0.8894

52%|██████ | 52/100 [09:05<08:22, 10.47s/it]

Training Loss : 0.11998906311511993, Testing Loss :
0.40435436005592346, Training Accuracy : 0.95704, Testing Accuracy :
0.8936

53%|██████ | 53/100 [09:15<08:12, 10.48s/it]

Training Loss : 0.11606632172584534, Testing Loss :
0.4392728645801544, Training Accuracy : 0.9597, Testing Accuracy :
0.8916

54%|██████ | 54/100 [09:26<08:01, 10.46s/it]

Training Loss : 0.10802373628616332, Testing Loss :
0.4287527338027954, Training Accuracy : 0.96098, Testing Accuracy :
0.89

55%|██████ | 55/100 [09:36<07:50, 10.45s/it]

Training Loss : 0.11218075738430024, Testing Loss :
0.4454464340686798, Training Accuracy : 0.96176, Testing Accuracy :
0.8899

56%|██████ | 56/100 [09:46<07:40, 10.47s/it]

Training Loss : 0.10551222920417785, Testing Loss :
0.40828085255622865, Training Accuracy : 0.9629, Testing Accuracy :
0.8933

57%|██████ | 57/100 [09:57<07:33, 10.54s/it]

Training Loss : 0.09989498874664307, Testing Loss :
0.4223347550392151, Training Accuracy : 0.964, Testing Accuracy :
0.8968

58%|██████ | 58/100 [10:08<07:23, 10.56s/it]

Training Loss : 0.09043127286434173, Testing Loss :
0.4521582398176193, Training Accuracy : 0.96802, Testing Accuracy :
0.8915

59%|██████ | 59/100 [10:18<07:10, 10.51s/it]

Training Loss : 0.08866517826795578, Testing Loss :
0.42655759048461916, Training Accuracy : 0.96712, Testing Accuracy :
0.8966

60%|██████ | 60/100 [10:29<07:00, 10.52s/it]

Training Loss : 0.08230563286304474, Testing Loss :
0.4526253000259399, Training Accuracy : 0.97016, Testing Accuracy :
0.8947

61%|██████ | 61/100 [10:39<06:50, 10.52s/it]

Training Loss : 0.08154015812158584, Testing Loss :
0.4504463098526001, Training Accuracy : 0.97072, Testing Accuracy :
0.8958

62%|██████████ | 62/100 [10:49<06:35, 10.40s/it]

Training Loss : 0.07949949363470077, Testing Loss :
0.4565501708030701, Training Accuracy : 0.97202, Testing Accuracy :
0.8959

63%|██████████ | 63/100 [11:00<06:25, 10.41s/it]

Training Loss : 0.0747082015299797, Testing Loss : 0.4420340425491333,
Training Accuracy : 0.97306, Testing Accuracy : 0.8995

64%|██████████ | 64/100 [11:11<06:19, 10.55s/it]

Training Loss : 0.07220506791591644, Testing Loss :
0.48928732318878176, Training Accuracy : 0.97422, Testing Accuracy :
0.8956

65%|██████████ | 65/100 [11:21<06:05, 10.45s/it]

Training Loss : 0.07177760385513306, Testing Loss :
0.42840618667602537, Training Accuracy : 0.97424, Testing Accuracy :
0.8998

66%|██████████ | 66/100 [11:31<05:52, 10.36s/it]

Training Loss : 0.06686747444868088, Testing Loss :
0.45723098306655885, Training Accuracy : 0.97728, Testing Accuracy :
0.9035

67%|██████████ | 67/100 [11:41<05:42, 10.39s/it]

Training Loss : 0.06354128739356994, Testing Loss : 0.472225291967392,
Training Accuracy : 0.97752, Testing Accuracy : 0.9008

68%|██████████ | 68/100 [11:52<05:34, 10.46s/it]

Training Loss : 0.06781714969158173, Testing Loss :
0.42682201583385465, Training Accuracy : 0.97536, Testing Accuracy :
0.9036

69%|██████████ | 69/100 [12:03<05:23, 10.45s/it]

Training Loss : 0.05613621583223343, Testing Loss :
0.47618203630447387, Training Accuracy : 0.9805, Testing Accuracy :
0.904

70%|██████████ | 70/100 [12:13<05:13, 10.46s/it]

Training Loss : 0.04830798606395721, Testing Loss :
0.42769108983278276, Training Accuracy : 0.98292, Testing Accuracy :
0.9111

71%|██████████ | 71/100 [12:23<04:59, 10.33s/it]

Training Loss : 0.04202377602219581, Testing Loss :
0.4650554542541504, Training Accuracy : 0.98532, Testing Accuracy :
0.9086

72%|██████████ | 72/100 [12:33<04:47, 10.25s/it]

Training Loss : 0.044704957706928256, Testing Loss :
0.4745908398151398, Training Accuracy : 0.9841, Testing Accuracy :
0.906

73%|██████████ | 73/100 [12:43<04:37, 10.28s/it]

Training Loss : 0.04317195358723402, Testing Loss :
0.47118928098678586, Training Accuracy : 0.98452, Testing Accuracy :
0.9054

74%|██████████ | 74/100 [12:54<04:30, 10.39s/it]

Training Loss : 0.038022401463091375, Testing Loss :
0.4802045246601105, Training Accuracy : 0.98612, Testing Accuracy :
0.907

75%|██████████ | 75/100 [13:05<04:21, 10.45s/it]

Training Loss : 0.038868962845504285, Testing Loss :
0.4588340826034546, Training Accuracy : 0.98592, Testing Accuracy :
0.9075

76%|██████████ | 76/100 [13:15<04:11, 10.47s/it]

Training Loss : 0.03363791984736919, Testing Loss :
0.4912823760032654, Training Accuracy : 0.98772, Testing Accuracy :
0.9074

77%|██████████ | 77/100 [13:26<04:02, 10.54s/it]

Training Loss : 0.03306935468971729, Testing Loss :
0.4800869453430176, Training Accuracy : 0.98842, Testing Accuracy :
0.9119

78%|██████████ | 78/100 [13:36<03:51, 10.54s/it]

Training Loss : 0.030943616832494737, Testing Loss :
0.5103559723854065, Training Accuracy : 0.98878, Testing Accuracy :
0.9073

79%|██████████ | 79/100 [13:47<03:41, 10.53s/it]

Training Loss : 0.02811643488228321, Testing Loss :
0.47872517070770265, Training Accuracy : 0.99016, Testing Accuracy :
0.9105

80%|██████████ | 80/100 [13:57<03:30, 10.51s/it]

Training Loss : 0.027139724824428557, Testing Loss :
0.5180482689857483, Training Accuracy : 0.99054, Testing Accuracy :
0.9116

81%|██████████ | 81/100 [14:08<03:18, 10.45s/it]

Training Loss : 0.026257721359729768, Testing Loss :
0.4880937006950378, Training Accuracy : 0.991, Testing Accuracy :
0.9117

82%|██████████ | 82/100 [14:18<03:06, 10.38s/it]

Training Loss : 0.022030702589005233, Testing Loss :
0.5105392691135406, Training Accuracy : 0.9921, Testing Accuracy :
0.9131

83%|██████████ | 83/100 [14:28<02:56, 10.40s/it]

Training Loss : 0.024398157700002192, Testing Loss :
0.5084094598770141, Training Accuracy : 0.99134, Testing Accuracy :
0.9082

84%|██████████ | 84/100 [14:39<02:47, 10.47s/it]

Training Loss : 0.024022529678046703, Testing Loss :
0.5260673250198364, Training Accuracy : 0.99164, Testing Accuracy :
0.9117

85%|██████████ | 85/100 [14:49<02:36, 10.44s/it]

Training Loss : 0.022933623305782677, Testing Loss :
0.5059761106491089, Training Accuracy : 0.99226, Testing Accuracy :
0.9112

86%|██████████ | 86/100 [15:00<02:26, 10.49s/it]

Training Loss : 0.021772636709790676, Testing Loss :
0.5110366032600403, Training Accuracy : 0.99248, Testing Accuracy :
0.9108

87%|██████████ | 87/100 [15:10<02:15, 10.45s/it]

Training Loss : 0.02018269987732172, Testing Loss :
0.5267858083724976, Training Accuracy : 0.99248, Testing Accuracy :
0.9106

88%|██████████ | 88/100 [15:21<02:05, 10.45s/it]

Training Loss : 0.020168822491690518, Testing Loss :
0.49368031721115113, Training Accuracy : 0.99354, Testing Accuracy :
0.918

89%|██████████ | 89/100 [15:31<01:55, 10.49s/it]

Training Loss : 0.013290383730456233, Testing Loss :
0.514165759563446, Training Accuracy : 0.9955, Testing Accuracy :
0.9159

90%|██████████ | 90/100 [15:42<01:45, 10.51s/it]

Training Loss : 0.01507324075654149, Testing Loss :
0.5301901911735535, Training Accuracy : 0.99432, Testing Accuracy :
0.9166

91%|██████████ | 91/100 [15:52<01:34, 10.52s/it]

Training Loss : 0.01492599542953074, Testing Loss :
0.5241970549583435, Training Accuracy : 0.995, Testing Accuracy :
0.9126

92%|██████████ | 92/100 [16:03<01:24, 10.54s/it]

Training Loss : 0.013875124659389258, Testing Loss :
0.516765375328064, Training Accuracy : 0.9952, Testing Accuracy :
0.918

93%|██████████ | 93/100 [16:13<01:13, 10.50s/it]

Training Loss : 0.01340749524973333, Testing Loss :
0.5306621287345886, Training Accuracy : 0.9955, Testing Accuracy :
0.9133

94%|██████████ | 94/100 [16:24<01:03, 10.55s/it]

Training Loss : 0.011001844610236585, Testing Loss :
0.5336188425540924, Training Accuracy : 0.9963, Testing Accuracy :
0.9159

95%|██████████ | 95/100 [16:34<00:52, 10.47s/it]

Training Loss : 0.013007677709460258, Testing Loss :
0.5420924195289611, Training Accuracy : 0.99538, Testing Accuracy :
0.9141

96%|██████████ | 96/100 [16:45<00:41, 10.45s/it]

Training Loss : 0.011571285721063614, Testing Loss :
0.5435201092720032, Training Accuracy : 0.99602, Testing Accuracy :
0.917

97%|██████████ | 97/100 [16:55<00:31, 10.44s/it]

Training Loss : 0.011715136720538139, Testing Loss :
0.5601223219871521, Training Accuracy : 0.99596, Testing Accuracy :
0.916

98%|██████████| 98/100 [17:06<00:20, 10.47s/it]

Training Loss : 0.011892944584768266, Testing Loss :
0.5583502286911011, Training Accuracy : 0.99608, Testing Accuracy :
0.9144

99%|██████████| 99/100 [17:16<00:10, 10.52s/it]

Training Loss : 0.012115279949195683, Testing Loss :
0.5437278823375702, Training Accuracy : 0.99578, Testing Accuracy :
0.9179

100%|██████████| 100/100 [17:27<00:00, 10.48s/it]

Training Loss : 0.008914166655614971, Testing Loss :
0.5569703786849975, Training Accuracy : 0.99684, Testing Accuracy :
0.9176

```
print("Maximum Testing Accuracy Achieved: %s"%(max(testingAccuracy)))  
xmax = np.argmax(testingAccuracy)  
ymax = max(testingAccuracy)
```

Maximum Testing Accuracy Achieved: 0.918

```
f, (ax1, ax2) = plt.subplots(1, 2, figsize = (20, 10))  
n = len(trainingLoss)  
ax1.plot(range(n), trainingLoss, '-', linewidth = '3', label = 'Train  
Error')  
ax1.plot(range(n), testingLoss, '-', linewidth = '3', label = 'Test  
Error')  
ax2.plot(range(n), trainingAccuracy, '-', linewidth = '3', label =  
'Train Accuracy')  
ax2.plot(range(n), testingAccuracy, '-', linewidth = '3', label =  
'Test Accuracy')  
ax2.annotate('max accuracy = %s'%(ymax), xy = (xmax, ymax), xytext =  
(xmax, ymax+0.15), arrowprops = dict(facecolor = 'black', shrink =  
0.05))  
ax1.grid(True)  
ax2.grid(True)  
ax1.legend()  
ax2.legend()  
f.savefig("./trainTestCurve.png")
```

